

Code Modification

```
## Makefile
```

```
(-)CS333_PROJECT ?= 1  
(+)CS333_PROJECT ?= 2
```

```
## uproc.h
```

```
(+)#include "types.h"
```

```
## defs.h
```

```
(+)#include "uproc.h"
```

```
// proc.c
```

```
int      cpuid(void);  
void     exit(void);  
int      fork(void);  
int      growproc(int);  
int      kill(int);  
struct cpu* mycpu(void);  
struct proc* myproc();  
void     pinit(void);  
void     procdump(void);  
void     scheduler(void) __attribute__((noreturn));  
void     sched(void);  
void     setproc(struct proc*);  
void     sleep(void*, struct spinlock*);  
void     userinit(void);  
int      wait(void);  
void     wakeup(void*);  
void     yield(void);  
(+)#ifdef CS333_P2  
(+)int      getprocs(uint max, struct uproc* upTable);  
(+)#endif  
#ifdef CS333_P3  
void     printFreeList(void);  
void     printList(int);  
void     printListStats(void);  
#endif // CS333_P3
```

```
## proc.c
```

```
(+)#ifdef CS333_P2  
(+) #include "uproc.h"  
(+)#endif
```

```
static struct proc*  
allocproc(void)  
{  
    struct proc *p;  
    char *sp;  
  
    acquire(&ptable.lock);  
    int found = 0;  
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)  
        if(p->state == UNUSED) {  
            found = 1;  
            break;  
        }  
    if (!found) {  
        release(&ptable.lock);  
    }
```

```

    return 0;
}
p->state = EMBRYO;
p->pid = nextpid++;
release(&ptable.lock);

// Allocate kernel stack.
if((p->kstack = kalloc()) == 0){
    p->state = UNUSED;
    return 0;
}
sp = p->kstack + KSTACKSIZE;

// Leave room for trap frame.
sp -= sizeof *p->tf;
p->tf = (struct trapframe*)sp;

// Set up new context to start executing at forkret,
// which returns to trapret.
sp -= 4;
*(uint*)sp = (uint)trapret;

sp -= sizeof *p->context;
p->context = (struct context*)sp;
memset(p->context, 0, sizeof *p->context);
p->context->eip = (uint)forkret;

#ifdef CS333_P1
    p->start_ticks = ticks;
#endif // CS333_P1

(+)#ifdef CS333_P2
(+)  p->cpu_ticks_total = 0;
(+)  p->cpu_ticks_in = 0;
(+)#endif // CS333_P2

return p;
}

void
userinit(void)
{
    struct proc *p;
    extern char _binary_initcode_start[], _binary_initcode_size[];

    p = allocproc();

    initproc = p;
    if((p->pgdir = setupkvm()) == 0)
        panic("userinit: out of memory?");
    inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
    p->sz = PGSIZE;
    memset(p->tf, 0, sizeof(*p->tf));
    p->tf->cs = (SEG_UCODE << 3) | DPL_USER;
    p->tf->ds = (SEG_UDATA << 3) | DPL_USER;
    p->tf->es = p->tf->ds;
    p->tf->ss = p->tf->ds;
    p->tf->eflags = FL_IF;
    p->tf->esp = PGSIZE;
    p->tf->eip = 0; // beginning of initcode.S

    (+)#ifdef CS333_P2
    (+)  p->uid = DEFAULT_UID;
    (+)  p->gid = DEFAULT_GID;

```

```

(+)#endif

safestrcpy(p->name, "initcode", sizeof(p->name));
p->cwd = namei("/");

// this assignment to p->state lets other cores
// run this process. the acquire forces the above
// writes to be visible, and the lock is also needed
// because the assignment might not be atomic.
acquire(&ptable.lock);
p->state = RUNNABLE;
release(&ptable.lock);
}

int
fork(void)
{
    int i;
    uint pid;
    struct proc *np;
    struct proc *curproc = myproc();

    // Allocate process.
    if((np = allocproc()) == 0){
        return -1;
    }

    // Copy process state from proc.
    if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0){
        kfree(np->kstack);
        np->kstack = 0;
        np->state = UNUSED;
        return -1;
    }
    np->sz = curproc->sz;
    np->parent = curproc;
    *np->tf = *curproc->tf;

    (+)#ifdef CS333_P2
    (+) np->uid = curproc->uid;
    (+) np->gid = curproc->gid;
    (+)#endif

    // Clear %eax so that fork returns 0 in the child.
    np->tf->eax = 0;

    for(i = 0; i < NOFILE; i++)
        if(curproc->ofile[i])
            np->ofile[i] = filedup(curproc->ofile[i]);
    np->cwd = idup(curproc->cwd);

    safestrcpy(np->name, curproc->name, sizeof(curproc->name));

    pid = np->pid;

    acquire(&ptable.lock);
    np->state = RUNNABLE;
    release(&ptable.lock);

    return pid;
}

void
scheduler(void)

```

```

{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;
    #ifdef PDX_XV6
        int idle; // for checking if processor is idle
    #endif // PDX_XV6

    for(;;){
        // Enable interrupts on this processor.
        sti();

        #ifdef PDX_XV6
            idle = 1; // assume idle unless we schedule a process
        #endif // PDX_XV6
        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            #ifdef PDX_XV6
                idle = 0; // not idle this timeslice
            #endif // PDX_XV6
            c->proc = p;
            switchvm(p);
            p->state = RUNNING;
            (+)#ifdef CS333_P2
            (+) p->cpu_ticks_in = ticks;
            (+)#endif // CS333_P2
            switch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);

        #ifdef PDX_XV6
            // if idle, wait for next interrupt
            if (idle) {
                sti();
                hlt();
            }
        #endif // PDX_XV6
    }
}

(+)#ifdef CS333_P2
(+)int
(+)getprocs(uint max, struct uproc* upTable){
(+) struct proc* p;
(+) int procsNumber = 0;
(+) acquire(&ptable.lock);
(+) for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
(+)     if (procsNumber < max) {
(+)         if(p->state != UNUSED && p->state != EMBRYO) {
(+)             if(p->state >= 0 && p->state < NELEM(states) && states[p->state]){
(+)                 safestrcpy(upTable[procsNumber].state, states[p->state], STRMAX);
(+)             } else {

```

```

(+)         safestrcpy(upTable[procsNumber].state, "???", STRMAX);
(+)     }
(+)     upTable[procsNumber].pid = p->pid;
(+)     upTable[procsNumber].uid = p->uid;
(+)     upTable[procsNumber].gid = p->gid;
(+)     upTable[procsNumber].ppid = p->parent ? p->parent->pid : p->pid;
(+)     upTable[procsNumber].elapsed_ticks = ticks - p->start_ticks;
(+)     upTable[procsNumber].CPU_total_ticks = p->cpu_ticks_total;
(+)     upTable[procsNumber].size = p->sz;
(+)     safestrcpy(upTable[procsNumber].name, p->name, STRMAX);
(+)     procsNumber++;
(+) }
(+) } else {
(+)     break;
(+) }
(+) }
(+) release(&ptable.lock);
(+) return procsNumber;
(+) }
(+) #endif // CS333_P2

```

proc.h

```

struct proc {
    uint sz;                // Size of process memory (bytes)
    pde_t* pgdir;           // Page table
    char *kstack;           // Bottom of kernel stack for this process
    enum procstate state;   // Process state
    uint pid;               // Process ID
    struct proc *parent;    // Parent process. NULL indicates no parent
    struct trapframe *tf;   // Trap frame for current syscall
    struct context *context; // swtch() here to run process
    void *chan;             // If non-zero, sleeping on chan
    int killed;             // If non-zero, have been killed
    struct file *ofile[NOFILE]; // Open files
    struct inode *cwd;       // Current directory
    char name[16];           // Process name (debugging)
    uint start_ticks;
    (+)uint uid;
    (+)uint gid;
    (+)uint cpu_ticks_total;
    (+)uint cpu_ticks_in;
};

```

ps.c

```

(+) #ifdef CS333_P2
(+) #include "types.h"
(+) #include "user.h"
(+) #include "uproc.h"

(+) #define MAX 16

(+) int
(+) main(void)
(+) {
(+)     struct uproc *proc = malloc(sizeof(struct uproc)*MAX);
(+)     int procsNumber = getprocs(MAX, proc);
(+)     printf(1, "PID\tName\t\tUID\tGID\tPPID\tElapsed\tCPU\tState\tSize\n");

(+)     int i;
(+)     for(i = 0; i<procsNumber; i++){
(+)         struct uproc currentProc = proc[i];
(+)         uint elapsedTicks = currentProc.elapsed_ticks;

```



```

extern int sys_wait(void);
extern int sys_write(void);
extern int sys_uptime(void);
#ifdef PDX_XV6
extern int sys_halt(void);
#endif // PDX_XV6
#ifdef CS333_P1
extern int sys_date(void);
#endif //CS333_P1
(+)#ifdef CS333_P2
(+)extern int sys_getuid(void);
(+)extern int sys_getgid(void);
(+)extern int sys_getppid(void);
(+)extern int sys_setuid(void);
(+)extern int sys_setgid(void);
(+)extern int sys_getprocs(void);
(+)#endif //CS333_P2

static int (*syscalls[])(void) = {
[SYS_fork]      sys_fork,
[SYS_exit]      sys_exit,
[SYS_wait]      sys_wait,
[SYS_pipe]      sys_pipe,
[SYS_read]      sys_read,
[SYS_kill]      sys_kill,
[SYS_exec]      sys_exec,
[SYS_fstat]     sys_fstat,
[SYS_chdir]     sys_chdir,
[SYS_dup]       sys_dup,
[SYS_getpid]    sys_getpid,
[SYS_sbrk]      sys_sbrk,
[SYS_sleep]     sys_sleep,
[SYS_uptime]    sys_uptime,
[SYS_open]      sys_open,
[SYS_write]     sys_write,
[SYS_mknod]     sys_mknod,
[SYS_unlink]    sys_unlink,
[SYS_link]      sys_link,
[SYS_mkdir]     sys_mkdir,
[SYS_close]     sys_close,
#ifdef PDX_XV6
[SYS_halt]      sys_halt,
#endif // PDX_XV6
#ifdef CS333_P1
[SYS_date]      sys_date,
#endif // CS333_P1
(+)#ifdef CS333_P2
(+) [SYS_getuid]  sys_getuid,
(+) [SYS_getgid]  sys_getgid,
(+) [SYS_getppid] sys_getppid,
(+) [SYS_setuid]  sys_setuid,
(+) [SYS_setgid]  sys_setgid,
(+) [SYS_getprocs] sys_getprocs,
(+)#endif // CS333_P2
};

## syscall.h

#define SYS_fork      1
#define SYS_exit      SYS_fork+1
#define SYS_wait      SYS_exit+1
#define SYS_pipe      SYS_wait+1
#define SYS_read      SYS_pipe+1
#define SYS_kill      SYS_read+1

```

```

#define SYS_exec      SYS_kill+1
#define SYS_fstat     SYS_exec+1
#define SYS_chdir     SYS_fstat+1
#define SYS_dup       SYS_chdir+1
#define SYS_getpid    SYS_dup+1
#define SYS_sbrk      SYS_getpid+1
#define SYS_sleep     SYS_sbrk+1
#define SYS_uptime    SYS_sleep+1
#define SYS_open      SYS_uptime+1
#define SYS_write     SYS_open+1
#define SYS_mknod     SYS_write+1
#define SYS_unlink    SYS_mknod+1
#define SYS_link      SYS_unlink+1
#define SYS_mkdir     SYS_link+1
#define SYS_close     SYS_mkdir+1
#define SYS_halt      SYS_close+1
#define SYS_date      SYS_halt+1
(+)#define SYS_getuid  SYS_date+1
(+)#define SYS_getgid  SYS_getuid+1
(+)#define SYS_getppid SYS_getgid+1
(+)#define SYS_setuid   SYS_getppid+1
(+)#define SYS_setgid   SYS_setuid+1
(+)#define SYS_getprocs SYS_setgid+1

```

testsetuid.c

```

(+)#ifdef CS333_P2
(+)#include "types.h"
(+)#include "user.h"

(+)int
(+)main(int argc, char *argv[])
(+) {
(+)    printf(1, "***** In %s: my uid is %d\n\n", argv[0], getuid());
(+)    exit();
(+)}
(+)#endif

```

user.h

```

// ulib.c
int stat(char*, struct stat*);
char* strcpy(char*, char*);
void *memmove(void*, void*, int);
char* strchr(const char*, char c);
int strcmp(const char*, const char*);
void printf(int, char*, ...);
char* gets(char*, int max);
uint strlen(char*);
void* memset(void*, int, uint);
void* malloc(uint);
void free(void*);
int atoi(const char*);
#ifdef PDX_XV6
int atoo(const char*);
int strncmp(const char*, const char*, uint);
#endif // PDX_XV6

(+)#ifdef CS333_P2
(+)uint getuid(void);
(+)uint getgid(void);
(+)uint getppid(void);

(+)int setuid(uint);

```



```
(+)int setgid(uint);

(+)int getprocs(uint max, struct uproc* table);
(+)#endif
```

```
## usys.S
```

```
SYSCALL(fork)
SYSCALL(exit)
SYSCALL(wait)
SYSCALL(pipe)
SYSCALL(read)
SYSCALL(write)
SYSCALL(close)
SYSCALL(kill)
SYSCALL(exec)
SYSCALL(open)
SYSCALL(mknod)
SYSCALL(unlink)
SYSCALL(fstat)
SYSCALL(link)
SYSCALL(mkdir)
SYSCALL(chdir)
SYSCALL(dup)
SYSCALL(getpid)
SYSCALL(sbrk)
SYSCALL(sleep)
SYSCALL(uptime)
SYSCALL(halt)
SYSCALL(date)
(+)SYSCALL(getuid)
(+)SYSCALL(getgid)
(+)SYSCALL(getppid)
(+)SYSCALL(setuid)
(+)SYSCALL(setgid)
(+)SYSCALL(getprocs)
```

```
## sysproc.c
```

```
#include "types.h"
#include "x86.h"
#include "defs.h"
#include "date.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#ifdef PDX_XV6
    #include "pdx-kernel.h"
#endif // PDX_XV6
```

```
(+)#ifdef CS333_P2
(+) #include "uproc.h"
(+)#endif // CS333_P2
```

```
(+)#ifdef CS333_P2
(+)int
(+)sys_getuid(void)
(+) {
(+)     struct proc *curproc = myproc();
(+)     return curproc->uid;
(+) }
```

```
(+)int
```

```

(+)sys_getgid(void)
(+) {
(+) struct proc *curproc = myproc();
(+) return curproc->gid;
(+) }

(+)int
(+)sys_getppid(void)
(+) {
(+) struct proc *curproc = myproc();
(+) struct proc *parent = curproc->parent;
(+) return parent != NULL ? parent->pid : 0;
(+) }

(+)int sys_setuid(void)
(+) {
(+) uint uid;
(+) struct proc *curproc = myproc();

(+) if(argint(0, (int*)&uid) >= 0) {
(+)     if(uid >= 0 && uid <= 32767) {
(+)         curproc->uid = uid;
(+)         return 0;
(+)     }
(+) }

(+) return -1;
(+) }

(+)int sys_setgid(void)
(+) {
(+) uint gid;
(+) struct proc *curproc = myproc();

(+) if(argint(0, (int*)&gid) >= 0) {
(+)     if(gid >= 0 && gid <= 32767) {
(+)         curproc->gid = gid;
(+)         return 0;
(+)     }
(+) }

(+) return -1;
(+) }

(+)int sys_getprocs(void)
(+) {
(+) uint max;
(+) struct uproc* proc;

(+) if (argint(0, (int*)&max) >= 0) {
(+)     if (max == 1 || max == 16 || max == 64 || max == 72) {
(+)         if (argptr(1, (void*)&proc, sizeof(struct uproc)) >= 0) {
(+)             return getprocs(max, proc);
(+)         }
(+)     }
(+) }

(+) return -1;
(+) }
(+)#endif //CS333_P2

```